

# pandas resample

November 26, 2018

ww# Examples of the use of *pandas.resample()* to calculate sum and average in different periods

**0.1 The input file has measurements of air temperature, relative humidity and radiation, as well as a timestamp field for each measurement. It includes one complete day of measurements, taken every 5 minutes, approximately. The measurements are *not* exactly equidistant in time (isochronal).**

**0.2 The function *pandas.resample()* will be used to help aggregate the measurements of temperature and radiation in two different ways:**

- \* The temperature will be averaged into hourly temperature
- \* The radiation will be integrated (added up)

```
#  
*
```

## 0.3 Import needed libraries

```
In [1]: %matplotlib inline  
import pandas as pd  
import matplotlib.pyplot as plt  
plt.ioff()
```

## 0.4 Read data

```
In [2]: df = pd.read_csv( '../ ../data/sensors_to_resample.csv', sep=';' )
```

### 0.4.1 First lines of the table, just to check

```
In [3]: df.head()
```

```
Out [3]:
```

	Timestamp	Temperature	Rel_Humidity	PAR_Radiation
0	2018-09-16 00:05:00	22.640	34.649	0.032
1	2018-09-16 00:10:00	22.699	35.280	0.024
2	2018-09-16 00:15:00	22.700	35.237	0.051
3	2018-09-16 00:20:00	22.700	34.618	0.013
4	2018-09-16 00:25:00	22.603	34.133	0.000

#### 0.4.2 And the last lines

```
In [4]: df.tail()
```

```
Out [4]:
```

		Timestamp	Temperature	Rel_Humidity	PAR_Radiation
282	2018-09-16	23:37:25	22.492	35.027	0.051
283	2018-09-16	23:42:25	22.386	34.442	0.051
284	2018-09-16	23:47:25	22.198	33.610	0.025
285	2018-09-16	23:52:26	22.172	35.041	0.025
286	2018-09-16	23:57:27	22.200	35.314	0.025

#### 0.5 We define the index of the table (previously a consecutive number) to be the data column *Timestamp*

```
In [5]: df.index = pd.DatetimeIndex( df['Timestamp'] )
df.index
```

```
Out [5]: DatetimeIndex(['2018-09-16 00:05:00', '2018-09-16 00:10:00',
                        '2018-09-16 00:15:00', '2018-09-16 00:20:00',
                        '2018-09-16 00:25:00', '2018-09-16 00:30:00',
                        '2018-09-16 00:35:01', '2018-09-16 00:40:02',
                        '2018-09-16 00:45:03', '2018-09-16 00:50:04',
                        ...,
                        '2018-09-16 23:12:23', '2018-09-16 23:17:24',
                        '2018-09-16 23:22:25', '2018-09-16 23:27:25',
                        '2018-09-16 23:32:25', '2018-09-16 23:37:25',
                        '2018-09-16 23:42:25', '2018-09-16 23:47:25',
                        '2018-09-16 23:52:26', '2018-09-16 23:57:27'],
                        dtype='datetime64[ns]', name='Timestamp', length=287, freq=None)
```

#### 0.6 Notice that the data type of the column was from "whatever" (*object*, in this case) into *DatetimeIndex* (datetime64 bits)

#### 0.7 This is the crucial step!

```
In [6]: print( df.index.dtype)
print( df['Timestamp'].dtype )
```

```
datetime64[ns]
object
```

#### 0.7.1 First lines of the table (you never check too much)

```
In [7]: df.head()
```

```
Out [7]:
```

		Timestamp	Temperature	Rel_Humidity	\
	Timestamp				
	2018-09-16 00:05:00	2018-09-16 00:05:00	22.640	34.649	
	2018-09-16 00:10:00	2018-09-16 00:10:00	22.699	35.280	

2018-09-16 00:15:00	2018-09-16 00:15:00	22.700	35.237
2018-09-16 00:20:00	2018-09-16 00:20:00	22.700	34.618
2018-09-16 00:25:00	2018-09-16 00:25:00	22.603	34.133

PAR_Radiation	
Timestamp	
2018-09-16 00:05:00	0.032
2018-09-16 00:10:00	0.024
2018-09-16 00:15:00	0.051
2018-09-16 00:20:00	0.013
2018-09-16 00:25:00	0.000

## 0.7.2 And the last lines again

```
In [8]: df.tail()
```

```
Out[8]:
```

	Timestamp	Temperature	Rel_Humidity	\
Timestamp				
2018-09-16 23:37:25	2018-09-16 23:37:25	22.492	35.027	
2018-09-16 23:42:25	2018-09-16 23:42:25	22.386	34.442	
2018-09-16 23:47:25	2018-09-16 23:47:25	22.198	33.610	
2018-09-16 23:52:26	2018-09-16 23:52:26	22.172	35.041	
2018-09-16 23:57:27	2018-09-16 23:57:27	22.200	35.314	

PAR_Radiation	
Timestamp	
2018-09-16 23:37:25	0.051
2018-09-16 23:42:25	0.051
2018-09-16 23:47:25	0.025
2018-09-16 23:52:26	0.025
2018-09-16 23:57:27	0.025

## 0.8 Notice that the timestamps are not exactly isochronal: that is our main problem here

```
In [9]: df['Timestamp'].head()
```

```
Out[9]: Timestamp
2018-09-16 00:05:00    2018-09-16 00:05:00
2018-09-16 00:10:00    2018-09-16 00:10:00
2018-09-16 00:15:00    2018-09-16 00:15:00
2018-09-16 00:20:00    2018-09-16 00:20:00
2018-09-16 00:25:00    2018-09-16 00:25:00
Name: Timestamp, dtype: object
```

```
In [10]: df['Timestamp'].tail()
```

```
Out[10]: Timestamp
2018-09-16 23:37:25    2018-09-16 23:37:25
```

```
2018-09-16 23:42:25    2018-09-16 23:42:25
2018-09-16 23:47:25    2018-09-16 23:47:25
2018-09-16 23:52:26    2018-09-16 23:52:26
2018-09-16 23:57:27    2018-09-16 23:57:27
Name: Timestamp, dtype: object
```

0.9 Just to show it more clearly, we can check the difference (discret derivative, for the sake of precision) between timestamps

```
In [11]: duration_between_timestamps = df.index.to_series().diff()
```

### 0.10 Convert to seconds to compare more easily


```
In [12]: duration_between_timestamps = duration_between_timestamps.dt.seconds
```

```
In [13]: print( duration_between_timestamps.min() )
          print( duration_between_timestamps.max() )
```

300.0  
302.0

### 0.11 Plot the number of seconds between measurements

```
In [14]: fig, ax = plt.subplots( nrows=1, ncols=1, figsize=(20,5) )
         ax.plot( df.index, duration_between_timestamps, linestyle='', marker='s' )
         plt.show()
```



0.11.1 Doesn't seem like a big deal, a shift of one or two seconds, but towards midnight it adds up to a couple of minutes, we don't have the measurements at the "minute 5" anymore

```
In [15]: df['Timestamp'].tail()
```

```
Out[15]: Timestamp
2018-09-16 23:37:25    2018-09-16 23:37:25
2018-09-16 23:42:25    2018-09-16 23:42:25
2018-09-16 23:47:25    2018-09-16 23:47:25
```

```

2018-09-16 23:52:26    2018-09-16 23:52:26
2018-09-16 23:57:27    2018-09-16 23:57:27
Name: Timestamp, dtype: object

```

## 0.12 This is the problem to solve, and we want to do it with *pandas.resample*

### 0.12.1 For *resample* to work properly, the series must have a time index, as stated in the [documentation](#):

Convenience method for frequency conversion and resampling of time series. Object must have a

(that was the reason to use *DatetimeIndex* previously)

## 0.13 If we have that index, we can do things like:

```
In [16]: print( df.index.min(), df.index.max() )
```

```
2018-09-16 00:05:00 2018-09-16 23:57:27
```

```
In [17]: print( df.resample( '1H' ).mean().index.min(), df.resample( '1H' ).mean().index.max() )
```

```
2018-09-16 00:00:00 2018-09-16 23:00:00
```

```
In [18]: print( df.resample( '5Min' ).mean().index.min(), df.resample( '5Min' ).mean().index.max() )
```

```
2018-09-16 00:05:00 2018-09-16 23:55:00
```

For the record, 5 minutes can be '5Min' or '5T' ('5M' are 5 months)

```
In [19]: print( df.resample( '5M' ).mean().index.min(), df.resample( '5M' ).mean().index.max() )
          print( df.resample( '5Min' ).mean().index.min(), df.resample( '5Min' ).mean().index.max() )
          print( df.resample( '5T' ).mean().index.min(), df.resample( '5T' ).mean().index.max() )
```

```
2018-09-30 00:00:00 2018-09-30 00:00:00
```

```
2018-09-16 00:05:00 2018-09-16 23:55:00
```

```
2018-09-16 00:05:00 2018-09-16 23:55:00
```

## 0.14 What happened?

### 0.14.1 What happened was:

```
.resample( 'new intervals' ).mean()
```

**0.14.2** The index is recalculated to 1 hour or 5 minutes, and the *mean(!)* is taken as the new value for the intervals. This is perhaps easier to see in the hourly example:

```
In [20]: df['Temperature'].head(15)
```

```
Out[20]: Timestamp
2018-09-16 00:05:00    22.640
2018-09-16 00:10:00    22.699
2018-09-16 00:15:00    22.700
2018-09-16 00:20:00    22.700
2018-09-16 00:25:00    22.603
2018-09-16 00:30:00    22.503
2018-09-16 00:35:01    22.403
2018-09-16 00:40:02    22.410
2018-09-16 00:45:03    22.429
2018-09-16 00:50:04    22.430
2018-09-16 00:55:05    22.400
2018-09-16 01:00:05    22.301
2018-09-16 01:05:05    22.191
2018-09-16 01:10:05    22.102
2018-09-16 01:15:06    22.110
Name: Temperature, dtype: float64
```

```
In [21]: df['Temperature'].resample( '1H' ).mean().head()
```

```
Out[21]: Timestamp
2018-09-16 00:00:00    22.537909
2018-09-16 01:00:00    22.256083
2018-09-16 02:00:00    22.310167
2018-09-16 03:00:00    22.398333
2018-09-16 04:00:00    22.500833
Freq: H, Name: Temperature, dtype: float64
```

**0.14.3** Other methods for resample are:

```
bfill()    # Backward fill
count()    # Number of values in the interval
ffill()    # Forward fill
first()    # Use the first (valid) data
last()     # Use the last (valid) data
max()      # Maximum value in the interval
mean()     # Mean of the interval
median()   # Median of values in the interval
min()      # Minimum value in the interval
nunique()  # Number of unique values
std()      # Standard deviation
sum()      # Sum of the values in the interval
var()      # Variance in the interval
```

#### 0.14.4 Here a complete list (still need to check it thoroughly, though)

```
In [22]: tmp = df['Temperature'].resample( '1H' )
        methods = [ method_name for method_name in dir(tmp) if callable(getattr(tmp, method_name)) ]
        methods = [ method_name for method_name in methods if not '_' in method_name ]
        print( methods )

['agg', 'aggregate', 'apply', 'asfreq', 'backfill', 'bfill', 'count', 'ffill', 'fillna', 'first', 'last', 'max', 'mean', 'median', 'min', 'ohlc', 'pct_change', 'resample', 'rolling', 'sample', 'shift', 'take', 'transform', 'unstack', 'var']
```

#### 0.15 A couple of examples

```
In [23]: df['Temperature'].head(15)
```

```
Out[23]: Timestamp
2018-09-16 00:05:00    22.640
2018-09-16 00:10:00    22.699
2018-09-16 00:15:00    22.700
2018-09-16 00:20:00    22.700
2018-09-16 00:25:00    22.603
2018-09-16 00:30:00    22.503
2018-09-16 00:35:01    22.403
2018-09-16 00:40:02    22.410
2018-09-16 00:45:03    22.429
2018-09-16 00:50:04    22.430
2018-09-16 00:55:05    22.400
2018-09-16 01:00:05    22.301
2018-09-16 01:05:05    22.191
2018-09-16 01:10:05    22.102
2018-09-16 01:15:06    22.110
Name: Temperature, dtype: float64
```

```
In [24]: df['Temperature'].resample( '1H' ).min().head(2)
```

```
Out[24]: Timestamp
2018-09-16 00:00:00    22.400
2018-09-16 01:00:00    22.102
Freq: H, Name: Temperature, dtype: float64
```

```
In [25]: df['Temperature'].resample( '1H' ).max().head(2)
```

```
Out[25]: Timestamp
2018-09-16 00:00:00    22.70
2018-09-16 01:00:00    22.41
Freq: H, Name: Temperature, dtype: float64
```

```
In [26]: df['Temperature'].resample( '1H' ).sum().head(2)
```

```
Out[26]: Timestamp
2018-09-16 00:00:00    247.917
2018-09-16 01:00:00    267.073
Freq: H, Name: Temperature, dtype: float64
```

```
In [27]: df['Temperature'].resample( '1H' ).first().head(2)
```

```
Out[27]: Timestamp
2018-09-16 00:00:00    22.640
2018-09-16 01:00:00    22.301
Freq: H, Name: Temperature, dtype: float64
```

## 0.16 Some of the methods make more sense for higher frequencies, when the resampling frequency is higher than the original:

```
In [28]: df['Temperature'].resample( '1H' ).ffill().head(2)
```

```
Out[28]: Timestamp
2018-09-16 00:00:00     NaN
2018-09-16 01:00:00    22.4
Freq: H, Name: Temperature, dtype: float64
```

```
In [29]: df['Temperature'].resample( '1H' ).bfill().head(2)
```

```
Out[29]: Timestamp
2018-09-16 00:00:00    22.640
2018-09-16 01:00:00    22.301
Freq: H, Name: Temperature, dtype: float64
```

### 0.16.1 Now the same with 1 minute as new period:

```
In [30]: df['Temperature'].resample( '1T' ).ffill().head(10)
```

```
Out[30]: Timestamp
2018-09-16 00:05:00    22.640
2018-09-16 00:06:00    22.640
2018-09-16 00:07:00    22.640
2018-09-16 00:08:00    22.640
2018-09-16 00:09:00    22.640
2018-09-16 00:10:00    22.699
2018-09-16 00:11:00    22.699
2018-09-16 00:12:00    22.699
2018-09-16 00:13:00    22.699
2018-09-16 00:14:00    22.699
Freq: T, Name: Temperature, dtype: float64
```

```
In [31]: df['Temperature'].resample( '1T' ).bfill().head(10)
```

```
Out[31]: Timestamp
2018-09-16 00:05:00    22.640
2018-09-16 00:06:00    22.699
2018-09-16 00:07:00    22.699
2018-09-16 00:08:00    22.699
2018-09-16 00:09:00    22.699
```



```

2018-09-16 00:10:00    22.699
2018-09-16 00:11:00    22.700
2018-09-16 00:12:00    22.700
2018-09-16 00:13:00    22.700
2018-09-16 00:14:00    22.700
Freq: T, Name: Temperature, dtype: float64

```

```
In [32]: df['Temperature'].resample( '1T' ).fillna('nearest').head(10)
```

```

Out[32]: Timestamp
2018-09-16 00:05:00    22.640
2018-09-16 00:06:00    22.640
2018-09-16 00:07:00    22.640
2018-09-16 00:08:00    22.699
2018-09-16 00:09:00    22.699
2018-09-16 00:10:00    22.699
2018-09-16 00:11:00    22.699
2018-09-16 00:12:00    22.699
2018-09-16 00:13:00    22.700
2018-09-16 00:14:00    22.700
Freq: T, Name: Temperature, dtype: float64

```

```
In [33]: df['Temperature'].resample( '1T' ).fillna('nearest', limit=1).head(10)
```

```

Out[33]: Timestamp
2018-09-16 00:05:00    22.640
2018-09-16 00:06:00    22.640
2018-09-16 00:07:00      NaN
2018-09-16 00:08:00      NaN
2018-09-16 00:09:00    22.699
2018-09-16 00:10:00    22.699
2018-09-16 00:11:00    22.699
2018-09-16 00:12:00      NaN
2018-09-16 00:13:00      NaN
2018-09-16 00:14:00    22.700
Freq: T, Name: Temperature, dtype: float64

```

**0.17 Lastly, if used on a complete DataFrame, it applies to all columns, so please check to see if that is what you actually want**

```
In [34]: df.resample('1H').mean().head()
```

```

Out[34]:
           Temperature  Rel_Humidity  PAR_Radiation
Timestamp
2018-09-16 00:00:00    22.537909      34.753909      0.016818
2018-09-16 01:00:00    22.256083      34.337333      0.060417
2018-09-16 02:00:00    22.310167      34.318000      0.016000
2018-09-16 03:00:00    22.398333      34.478500      0.016917
2018-09-16 04:00:00    22.500833      34.427417      0.067667

```